




FIN GENIE: A Gen AI- Powered Financial Reports Tool

**Simplifying report summaries to facilitate
informed financial decisions**



Our Team



Soham Agarwal
Purdue MS-BAIM



Durga Dash
Purdue MS-BAIM



Anto Frederic
Purdue MS-BAIM



Chaitanya Varma
Purdue MS-BAIM



Bheeshma Ramaraju
Purdue MS-BAIM

*"The essence of a triumphant tale lies not in its plot, but in the hearts
and minds of those who bring it to life"*



Agenda



Future Scope

A lot of ideas we would love to implement

06

Results and Learnings

Our success stories and setbacks-turned-lessons

05

Product Demo

The fun part!

04

01

Business problem and benefits

Why does it really matter?

Approach

"80% of the work comes from the approach (the 20%)"
Durga

02

03

Phase 1 and Phase 2

Diving deeper into the sprints and phases

Business Problem

The Issue

Financial investors need additional information to make decisions. Thus additional sources

The deeper issue

SEC filings provide detailed audited reports. Reading 100 page documents isn't easy!

Existing processes...

Are Time intensive

Sorting through SEC data for additional information required considerable time and resources to gain insights.



Have High Opportunity costs

The opportunity cost is high in the fast-paced finance industry, making traditional analysis methods inefficient.



Business Benefits

Expected hours saved

At least 50% increase
in efficiency

Existing Methods

Days



Our Method

Hours



Increased efficiency

Automation reduces time spent on report generation, to give more time for investors to analyze information.



Cost Reduction

Less manual analysis means more hours saved and better investment decisions.



Scalability

Ability to handle increased data volumes efficiently as the number of documents grows.



User engagement

Interactive AI assistance enhances user experience and allows two-way engagement with the financial data.

The Timeline

Dates: January 12th 2024 – April 2nd 2024



January

WIREFRAMING

Designing the process flow

February

PHASE 1

Experimenting with existing tools

Feb + March

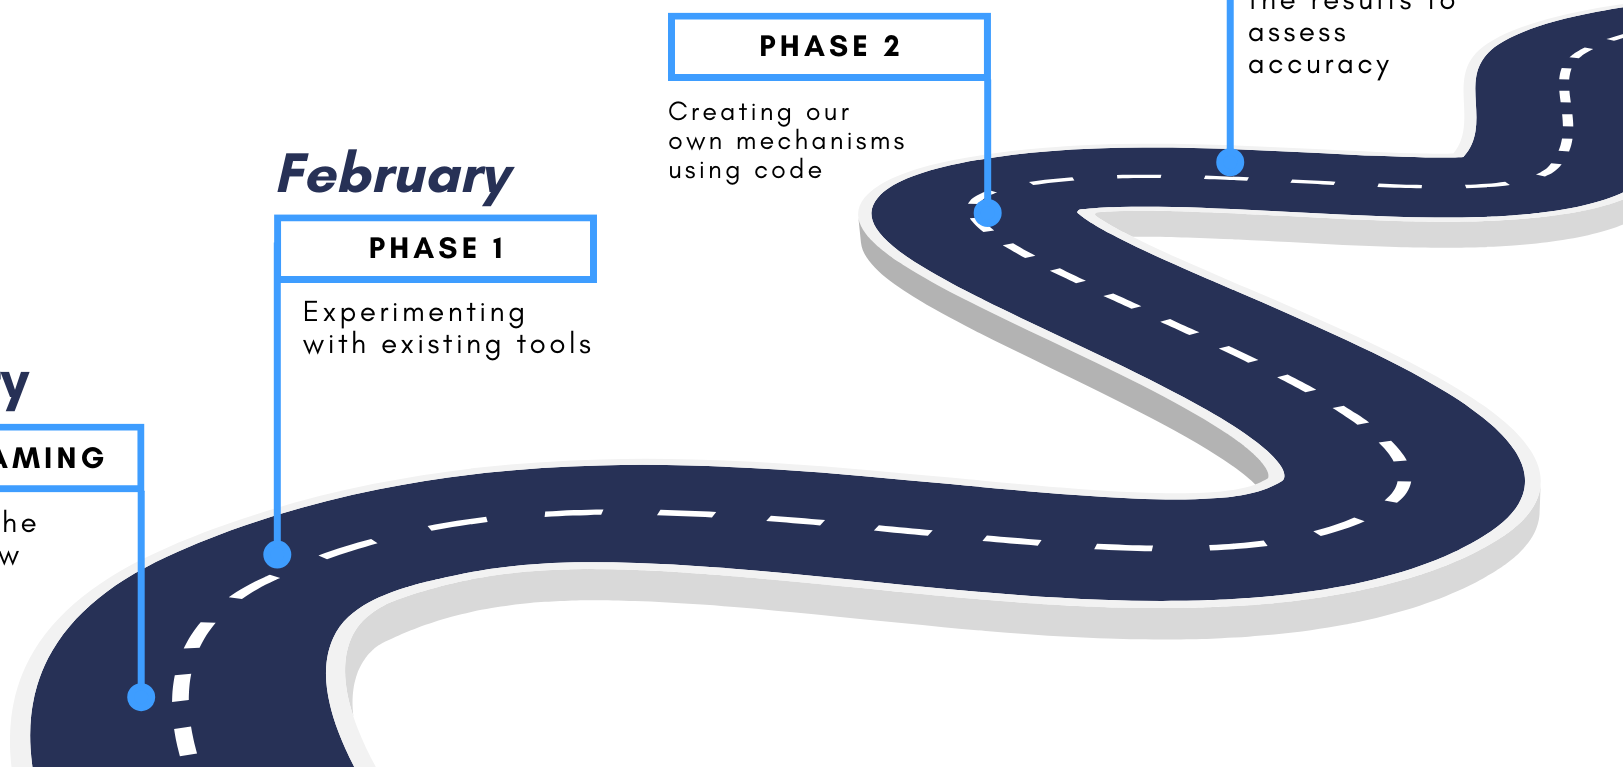
PHASE 2

Creating our own mechanisms using code

March

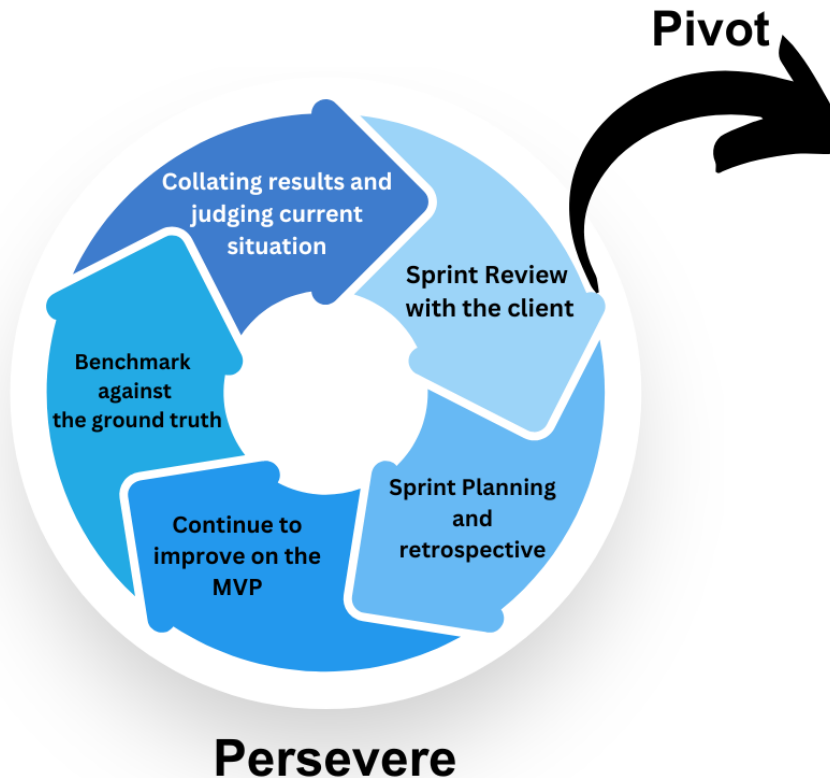
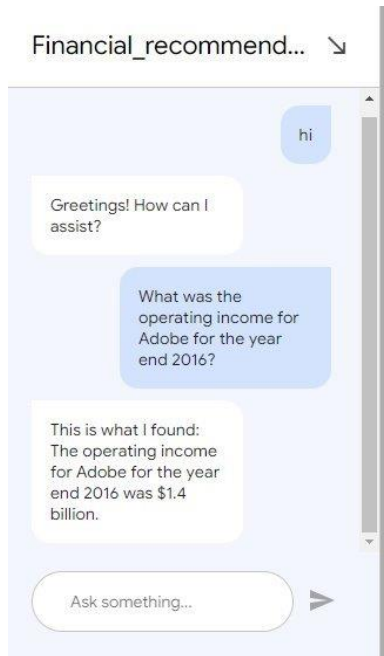
BENCHMARKING

Benchmarking the results to assess accuracy




Agile Approach

1) Older model



2) Newer model



Phase 1 Implementations

Step One -

Scraping 10k documents from the SEC website.




Step Three -

Using the large language models on GCP via the Search and Conversation, and Dialogflow platform.. Powered by Vertex AI.



Step Two -

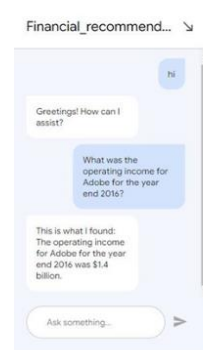
Storing the documents in the cloud to ensure scalability.



| | | |
|--------------------------|----------------------------------|----------|
| <input type="checkbox"/> | Adobe_2015.pdf | 1.4 MB |
| <input type="checkbox"/> | Adobe_2016.pdf | 4 MB |
| <input type="checkbox"/> | Adobe_2017.pdf | 3.8 MB |
| <input type="checkbox"/> | Amazon_2017.pdf | 612.4 KB |
| <input type="checkbox"/> | Amazon_2019.pdf | 604.7 KB |
| <input type="checkbox"/> | BESTBUY_2017.pdf | 758.2 KB |
| <input type="checkbox"/> | BESTBUY_2019.pdf | 763.8 KB |
| <input type="checkbox"/> | BLOCK_2016.pdf | 844.1 KB |
| <input type="checkbox"/> | BLOCK 2020.pdf | 1.1 MB |

Step Four -

Designing a front end chatbot and integrating this with our hyperparameter tuned LLMs



Wireframed Pipeline – Phase 1



Data Extraction

A combination of multiple extraction techniques to get accurate data.



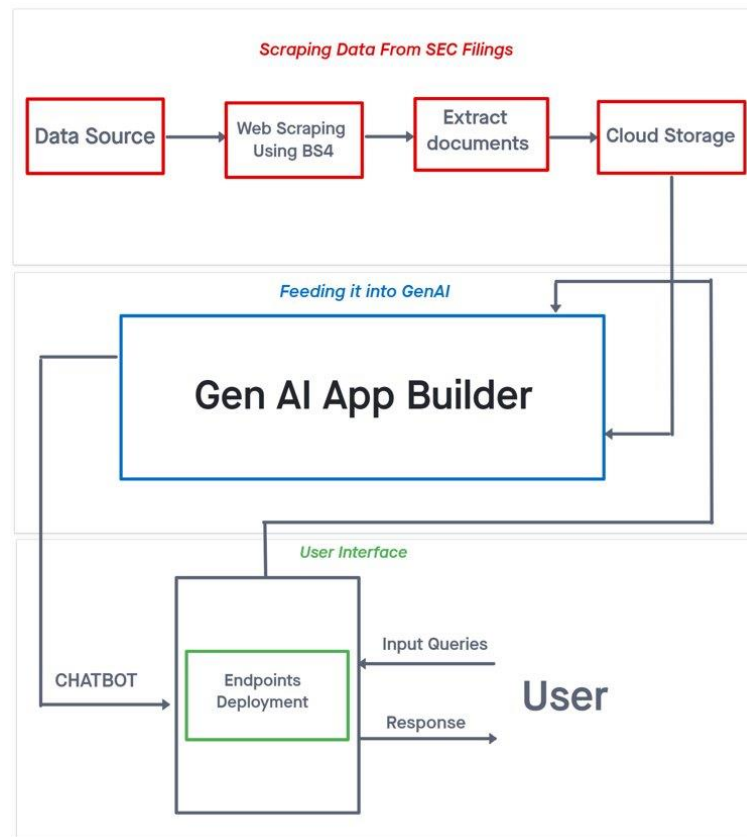
Data Interpretation

Google's cloud service "Search and Conversation" allows for an 'out-of-box' solution.



Data Delivery

Allows the user to have conversations with the data.



Phase 2 Implementations

Objective

- More control over responses generated.
- Ability to tailor responses to the financial domain
- Make a solution that shows an improvement over the Search & Conversation cloud service.



Additional Step 1 -

- Preprocessing SEC filings in batches with techniques like tokenization and lemmatization and chunking for context.
- vector stores are created using these chunks for efficient retrieval.



*Web Scraping and cloud storage methods were the same as in phase 1.

Phase 2 Implementations

Additional Step 2 -



- API's are called within a code environment. Prompting techniques, Query rephrasing, Chunk size strategies, few-shot prompting and many other techniques were tried to give us superior results.
- Models used: 1) Gecko embeddings + TextBison LLM model 2) Text Bison model.



Additional Step 3 -



- The top matching chunks are extracted to provide context in the prompt. The techniques applied above help for a more accurate retrieval of data.
- LLM generates the financial information and sends it back to the investor.

Wireframed Pipeline – Phase 2



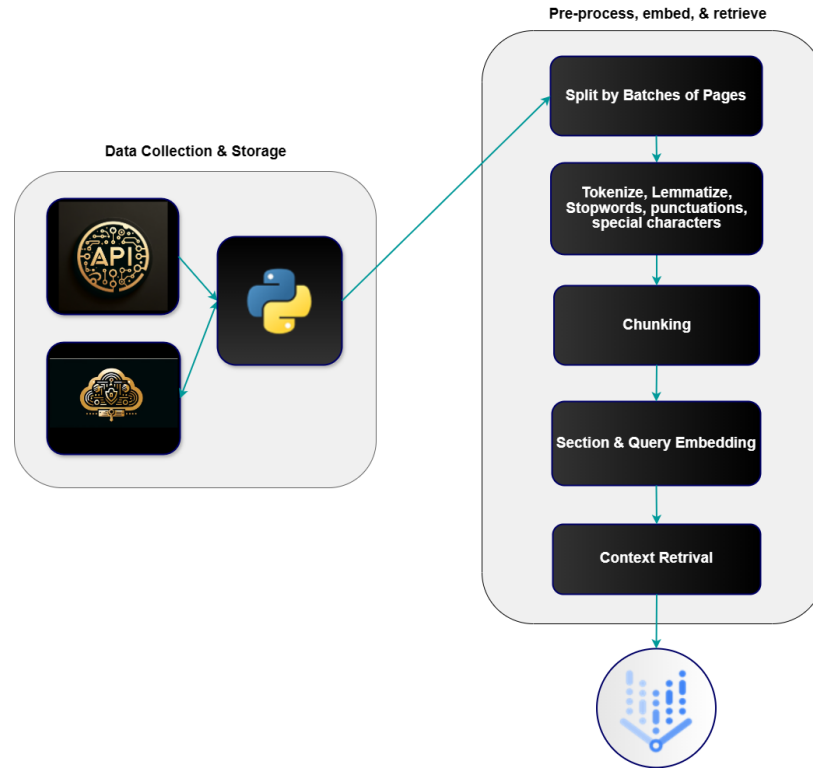
Preprocessing

Building a framework to allow model to fetch relevant information.



Chunking, embedding and retrieval

Helping the LLM understand how the document is structured.



Wireframed Pipeline – Phase 2



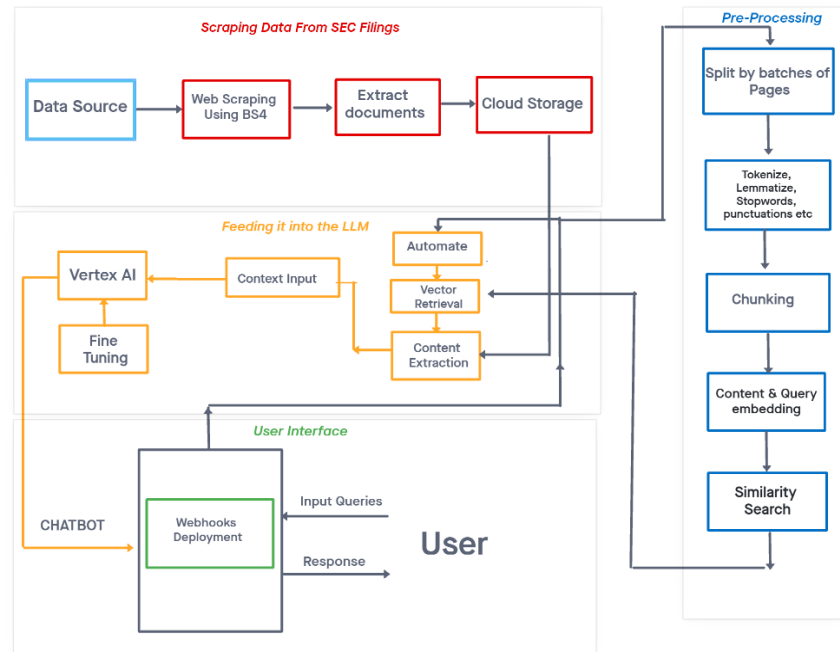
Cloud first

Using Google Cloud Platform (GCP) to stitch together technical components.

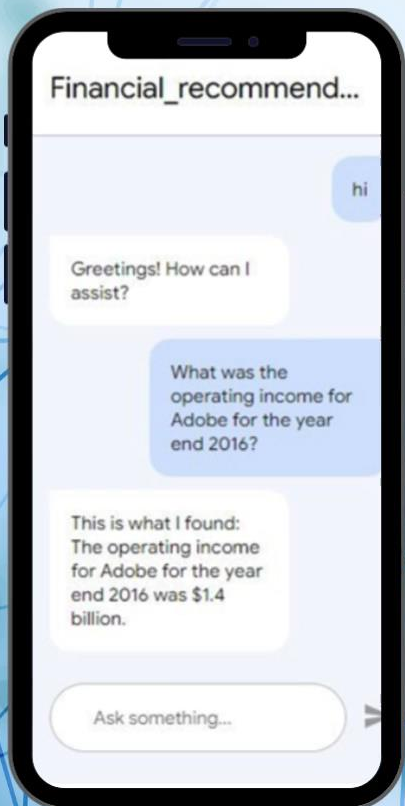


Feeding the LLM and hyperparameter tuning

Using the latest practices in GenAI development to enhance responses delivered in the finance domain.



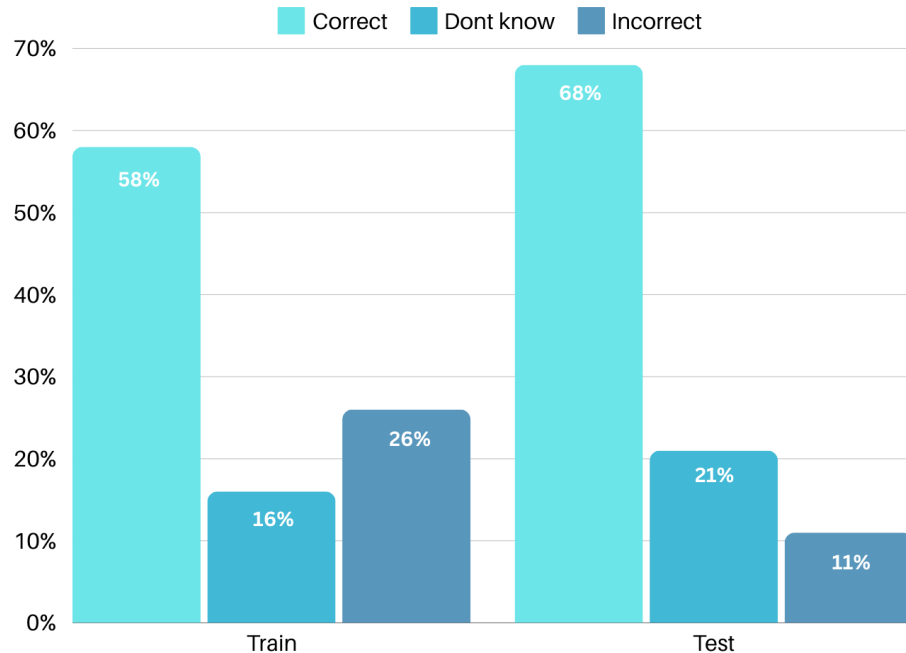
Product Demo



Results



Best Model Performance



Test vs Train




- Our model demonstrated comparable performance across both training and testing phases
- Noticeable decrease in the percentage of incorrect outcomes alongside these improvements
- Indicates a robust and reliable performance


Results



Model Performance



Model Performances



- The transition from Vertex AI to the TextBison base model resulted in a **2X** performance boost
- Followed by a significant **6X** enhancement when moving from the TextBison base to the TextBison final model

Learnings

What worked?

- Query Rephrasing.
- Page Batches/ Chunk Size.
- Data cleaning.
- Increased context window.
- One Shot prompting.

What didn't work?

- Overlap of content in chunking.
- Tokenizing and increasing prompt similarity.
- Using the ReAct framework.
- Fine-tuning the model on example questions.
- Cascading model Bert+Bison did not yield significant change.

Additional Learnings

Importance of model selection

Developed a model leveraging BART
Large CNN to condense 10K documents
into structured summaries by sections

Balancing Brevity and Detail

Fine-tuning summary parameters
underscores the importance of striking
a balance between brevity and detail

Document Summarization

Summary of Section 2: We earn tradable credits in the operation of our business under various regulations relating to greenhouse gas, fuel economy and clean fuel. We sell these credits to other regulated entities who can use the credits to comply with regulatory requirements. Sales of these credits are recognized within our consolidated financial statements included elsewhere in this Annual Report on Form 10-K.

Summary of Section 3: We may be unable to meet our projected construction timelines, costs and production ramp-ups, or we may experience difficulties in generating and maintaining demand for products manufactured there. All depend on our ability to continue to expand our sales capabilities. If we experience production delays or if demand recast, our business, financial condition and operating results may be harmed.

Summary of Section 4: Tesla's future growth and success are dependent upon consumers' demand for electric vehicles. The prices and availability of raw materials such as lithium, nickel, cobalt and/or other materials are unstable, depending on market conditions and global demand.

Summary of Section 5: The target demographics for our vehicles, particularly Model 3 and Model Y, are highly cyclical. Sales of vehicles tend to be cyclical in many markets, which may expose us to further volatility. Increased competition may result in lower vehicle unit sales, price reductions, revenue shortfalls, loss of customers and loss of market share. Handling of these products may cause disruption to the operation of our facilities.

Summary of Section 6: As of December 31, 2022, we and our subsidiaries had outstanding \$2.06 billion in aggregate amount of debt. Our ability to make scheduled payments of the principal and interest on our indebtedness depends on our future performance, which is subject to economic, financial, competitive and other factors beyond our control.

Summary of Section 7: We are exposed to fluctuations in currency exchange rates. We may need to defend ourselves against intellectual property infringement claims, which may be time-consuming and expensive. There can be no assurance that we will be able to adequately identify and protect the portions of intellectual property that are strategic to our business. We may be subject to potential suits or other legal demands by our competitors.

Additional Learnings

BERT Model...



PDF Embedding Processing

Read documents in batches, extract text per section, preprocess and generate semantic embeddings.



Context Retrieval

For a given question, the most relevant document sections are identified using embedding similarities.



Vector Store Creation

A store (dataframe) of texts and their embeddings is created to facilitate similarity searches.



Question Answering

With the context narrowed down, the question is answered using the fine-tuned BERT QA model

What new was discovered with this approach is that the context was retrieved very well without much pre-processing using the BERT model, but the model has troubled response generation which needs to be improved in the future scope of work. As part of this, research on how to tune context window, domain specific fine tuning, hugging face pipelines are being explored.

Future Scope



A combination of models to accurately identify tabular data and avoid misinterpretation.

Use pre-defined examples for model fine-tuning to enhance accuracy.

Create a benchmarking framework for comprehensive financial data analysis

Pair LLMs with mathematical libraries to improve accuracy in answering quantitative questions.

Develop an embedding model on extensive financial data to enhance context comprehension and information extraction.

Thank you for your support!



Ashwin Mishra



Prof. Matthew Lanham



Purdue

Thank You!

We will now take any questions.



Appendix

Appendix 1

UNITED STATES
SECURITIES AND EXCHANGE COMMISSION
Washington, D.C. 20549

FORM 10-K

☒ **ANNUAL REPORT PURSUANT TO SECTION 13 OR 15(d) OF THE SECURITIES EXCHANGE ACT OF 1934**

For the Fiscal Year Ended June 30, 2022

OR

☐ **TRANSITION REPORT PURSUANT TO SECTION 13 OR 15(d) OF THE SECURITIES EXCHANGE ACT OF 1934**

For the Transition Period From _____ to _____

Commission File Number 001-37845

MICROSOFT CORPORATION

WASHINGTON
(STATE OF INCORPORATION)

91-1144442
(I.R.S. ID)

ONE MICROSOFT WAY, REDMOND, WASHINGTON 98052-6399

(425) 882-8080

www.microsoft.com/investor

Securities registered pursuant to Section 12(b) of the Act:

Title of each class

Common stock, \$0.00000625 par value per share

3.125% Notes due 2028

2.625% Notes due 2033

Trading Symbol

MSFT

MSFT

MSFT

Name of exchange on which registered

NASDAQ

NASDAQ

NASDAQ

Securities registered pursuant to Section 12(g) of the Act:

Appendix 2

| | | <i>Page</i> |
|-----------------|---|-------------|
| PART I | | |
| Item 1. | Business | 3 |
| | Information about our Executive Officers | 21 |
| Item 1A. | Risk Factors | 23 |
| Item 1B. | Unresolved Staff Comments | 37 |
| Item 2. | Properties | 37 |
| Item 3. | Legal Proceedings | 37 |
| Item 4. | Mine Safety Disclosures | 37 |
| PART II | | |
| Item 5. | Market for Registrant's Common Equity, Related Stockholder Matters, and Issuer Purchases of Equity Securities | 38 |
| Item 6. | [Reserved] | 39 |
| Item 7. | Management's Discussion and Analysis of Financial Condition and Results of Operations | 40 |
| Item 7A. | Quantitative and Qualitative Disclosures about Market Risk | 56 |
| Item 8. | Financial Statements and Supplementary Data | 57 |
| Item 9. | Changes in and Disagreements with Accountants on Accounting and Financial Disclosure | 99 |
| Item 9A. | Controls and Procedures | 99 |
| | Report of Management on Internal Control over Financial Reporting | 99 |
| | Report of Independent Registered Public Accounting Firm | 100 |
| Item 9B. | Other Information | 101 |
| Item 9C. | Disclosure Regarding Foreign Jurisdictions that Prevent Inspections | 101 |
| PART III | | |
| Item 10. | Directors, Executive Officers and Corporate Governance | 101 |
| Item 11. | Executive Compensation | 101 |
| Item 12. | Security Ownership of Certain Beneficial Owners and Management and Related Stockholder Matters | 101 |
| Item 13. | Certain Relationships and Related Transactions, and Director Independence | 101 |
| Item 14. | Principal Accountant Fees and Services | 101 |
| PART IV | | |
| Item 15. | Exhibit and Financial Statement Schedules | 102 |
| Item 16. | Form 10-K Summary | 108 |

Appendix 3


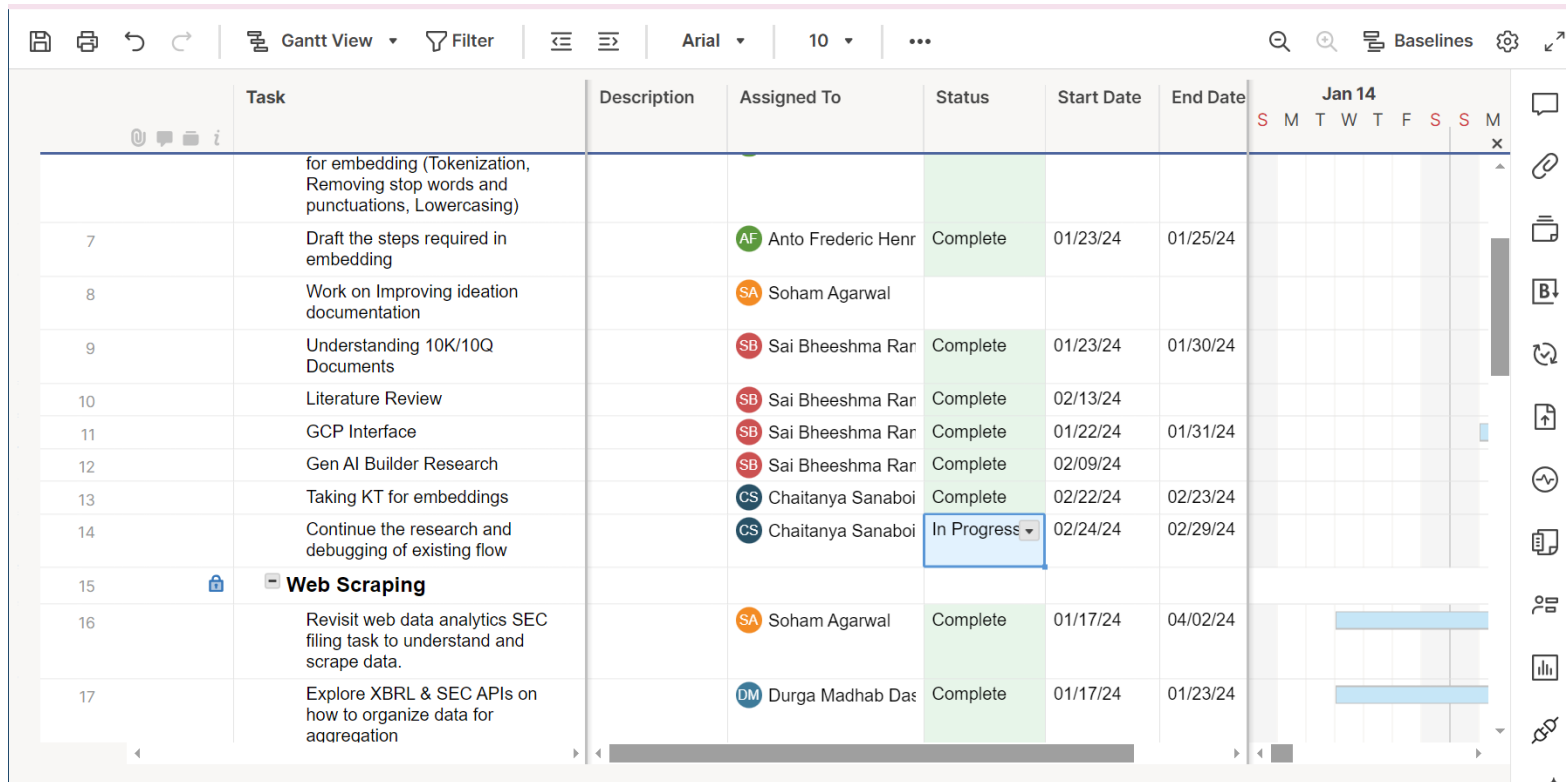


Fig. 3. Trend over time in the proportion of 10-K filings containing at least one infographic or other image. The sample consists of 47,906 10-K filings from 2003-2019. We define all variables in Appendix C and provide additional information on infographic types in Appendix A. The full-color version is available online.

Appendix 4



Appendix 5

```
embeddings_model = TextEmbeddingModel.from_pretrained("textembedding-gecko@001")
file_text_model = TextGenerationModel.from_pretrained("text-bison@001")
text_model = TextGenerationModel.from_pretrained("text-bison-32k")
rephrase_model = TextGenerationModel.from_pretrained("text-unicorn@001")

chat_model = ChatModel.from_pretrained("chat-bison-32k")
chat = chat_model.start_chat(
    context="You are an expert financial analyst. You know how to read SEC filings such as 10-K
")

gemini_pro_model = GenerativeModel("gemini-1.0-pro")
```

Models Tested & Used

Appendix 6

```
rephrase_prompt=f"""Act as a financial analyst with expert knowledge of SEC 10-K documents.  
Could you rephrase this question so that a financial novice can easily understand - {{USER_QUERY}}  
USER_QUERY: ``{question}``\n  
Please make sure to re-phrase the query such that it is 100% same in meaning to the original query  
"""
```

```
user_query = rephrase_model.predict(rephrase_prompt, temperature=0, top_k=1, max_output_tokens=300)  
user_query = user_query.replace('\n', ' ')  
user_query = add_fy_to_years(user_query)
```

```
lemmatized_query = []
```

```
for word in word_tokenize(user_query):  
    lemmatized_query.append(lemmatizer.lemmatize(word))
```

```
user_query_proc = " ".join([word.lower() for word in lemmatized_query if word not in string.punctuation])
```

Pre-Processing (User-Query)

Appendix 7 Pre-Processing (PDF Content)

```
section_content = read_pages_in_batches(file_path_pdf, 6)
```

```
def read_pages_in_batches(file_path_pdf, pages_per_batch=2):
    section_content = []

    with open(file_path_pdf, 'rb') as pdf_file:
        pdf_reader = PyPDF2.PdfReader(pdf_file)
        total_pages = len(pdf_reader.pages)
        #doc = fitz.open(pdf_file)
        #total_pages = len(doc)

        for page_num in range(0, total_pages, pages_per_batch):
            text = ""
            for i in range(pages_per_batch):
                if page_num + i < total_pages:
                    text += pdf_reader.pages[page_num + i].extract_text()
                    #text += doc.load_page(page_num + i).get_text()

            section_content.append(text)

    return section_content
```

Appendix 8

```
# Tokenize and Lemmatize the content in section_content
```

```
lemmatized_content = []
```

```
for section in section_content:
```

```
    lemmatized_section = []
```

```
    for word in word_tokenize(section):
```

```
        lemmatized_section.append(lemmatizer.lemmatize(word))
```

```
    lemmatized_content.append(" ".join(lemmatized_section))
```

```
preprocessed_sections = []
```

```
for content in lemmatized_content:
```

```
    proc_sections = [ word for word in content.lower().split() if word not in string.punctuation]
```

```
    preprocessed_sections.append(' '.join(proc_sections))
```

Appendix 9

Create Vector Store

```
text_lengths = [len(t) for t in preprocessed_sections]
size_overlap = int(np.median(text_lengths)//2)

CHUNK_SIZE = size_overlap
OVERLAP = 500
preprocessed_sections = preprocessed_sections[:len(preprocessed_sections)-1]

vector_store = create_vector_store(preprocessed_sections, CHUNK_SIZE, OVERLAP)
vector_store.to_pickle(file_path_pkl)
```

```
def create_vector_store(texts, chunk_size, overlap):
    vector_store = pd.DataFrame()

    # Split texts and filter out empty ones during splitting
    non_empty_texts = [text for text in itertools.chain.from_iterable(
        split_overlap(t, chunk_size, overlap) for t in texts if t.strip())]

    vector_store["texts"] = non_empty_texts

    # Create embeddings from those texts
    vector_store["embeddings"] = (
        vector_store["texts"].progress_apply(get_embeddings).apply(np.array)
    )

    return vector_store
```


Appendix 10

```
# Separates seq into multiple chunks in the specified size with the specified overlap
def split_overlap(seq, size, overlap):
    if len(seq) <= size:
        return [seq]
    return ["".join(x) for x in zip(*[seq[i :: size - overlap] for i in range(size)])]
```

```
@retry.Retry(timeout=300.0)
def get_embeddings(text):
    return embeddings_model.get_embeddings([text])[0].values
```

Appendix 11

Retrieving Context

```
def answer_question(user_query_proc, question, vector_store, num_docs=12, print_prompt=False):  
    context = get_context(user_query_proc, vector_store, num_docs)
```

```
def get_context(question, vector_store, num_docs):  
    # Tokenize the user query  
    query_tokens = word_tokenize(question)  
  
    # Embed the search query  
    query_vector = np.array(get_embeddings(" ".join(query_tokens)))  
  
    query_vector = np.array(get_embeddings(question))  
  
    # Get similarity to all other vectors and sort, cut off at num_docs  
    top_matched = (  
        vector_store["embeddings"]  
        .apply(get_similarity_fn(query_vector))  
        .sort_values(ascending=False)[:num_docs]  
        .index  
    )  
    top_matched_df = vector_store[vector_store.index.isin(top_matched)][["texts"]]  
  
    # Return a string with the top matches  
    context = " ".join(top_matched_df.texts.values)  
    #context = top_matched_df.texts.values  
    return context
```

Appendix 12

```
# Compute the cosine similarity of two vectors, wrap as returned function to make easier to use with
def get_similarity_fn(query_vector):
    def fn(row):
        return np.dot(row, query_vector) / (
            numpy.linalg.norm(row) * numpy.linalg.norm(query_vector)
        )

    return fn
```

Appendix 13

One-Shot Prompting for result

```
def answer_question(user_query_proc, question, vector_store, num_docs=12, print_prompt=False):
    context = get_context(user_query_proc, vector_store, num_docs)

    while len(context) < 100000:
        num_docs += 1
        context = get_context(user_query_proc, vector_store, num_docs)

    print(len(context))

    train_context = "Table of Contents ACTIVISION BLIZZARD, INC. AND SUBSIDIARIES CONSOLIDATED S
    train_question = "What is the FY2017 - FY2019 3 year average of capex as a % of revenue for
    train_answer = "To calculate the three-year average of capital expenditures (capex) as a per

    qa_prompt = f"""Input Example: Your mission is to answer questions based on a given context
    Context: ```{train_context}```
    Question: ***{train_question}***
    Before you give an answer, make sure it is only from information in the context. If the info
    Reply concisely.
    Answer: ```{train_answer}```\n
    Your mission is to answer questions based on a given context. Remember that before you give
    Context: ```{context}```
    Question: ***{question}***
    Before you give an answer, make sure it is only from information in the context. If the info
    Reply concisely.
    Answer: """

    if print_prompt:
        print(qa_prompt)
    result = text_model.predict(qa_prompt, temperature=0, top_k=3, max_output_tokens=400)
```